

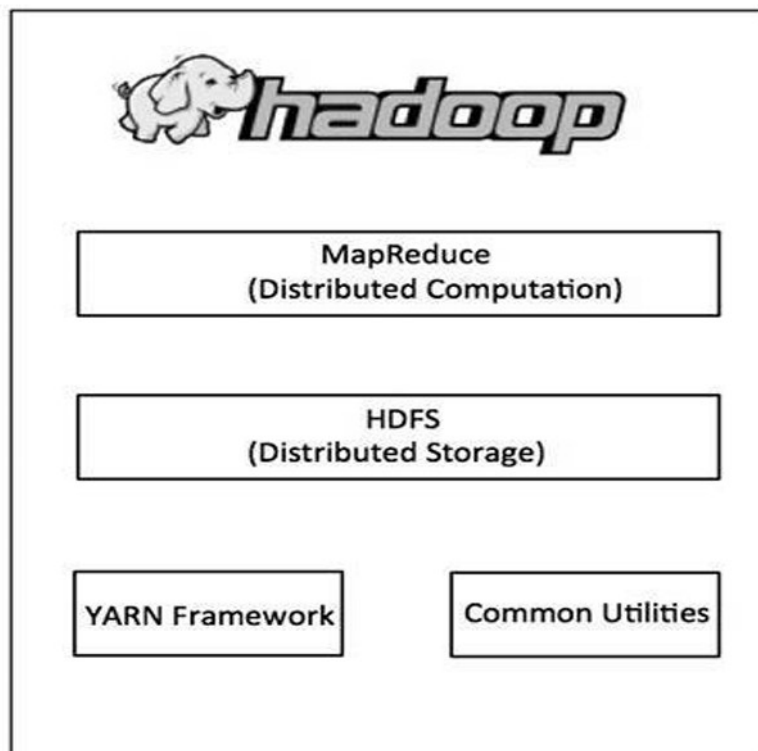
## **HADOOP**

*Hadoop is an Apache open source framework written in java that allows distributed processing of large datasets across clusters of computers using simple programming models. The Hadoop framework application works in an environment that provides distributed storage and computation across clusters of computers. Hadoop is designed to scale up from single server to thousands of machines, each offering local computation and storage.*

### **HADOOP ARCHITECTURE**

*At its core, Hadoop has two major layers namely:*

- (a) Processing/Computation layer (MapReduce).*
- (b) Storage layer (Hadoop Distributed File System).*



## Hadoop Distributed File System

The Hadoop Distributed File System (HDFS) is based on the Google File System (GFS) and provides a distributed file system that is designed to run on commodity hardware.

It has many similarities with existing distributed file systems. However, the differences from other distributed file systems are significant. It is highly fault-tolerant and is designed to be deployed on low-cost hardware.

It provides high throughput access to application data and is suitable for applications having large datasets. Apart from the above-mentioned two core components,

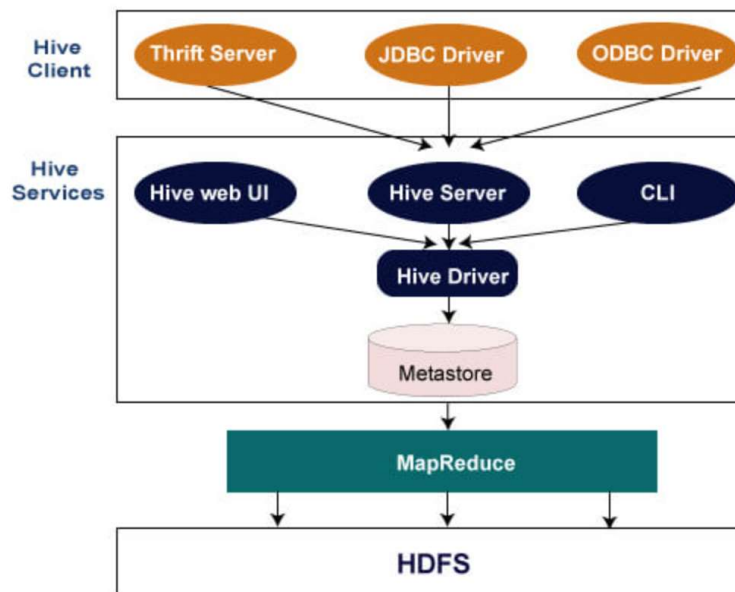
Hadoop framework also includes the following two modules:

- **Hadoop Common:** These are Java libraries and utilities required by other Hadoop modules.
- **Hadoop YARN:** This is a framework for job scheduling and cluster resource management.

## HIVE :

The Apache Hive is a distributed, fault-tolerant data warehouse system that enables analytics at a massive scale and facilitates reading, writing, and managing petabytes of data residing in distributed storage using SQL.

## Hive Architecture



**To Check the Database inside the HDFS :**

```
hive> show databases;  
OK  
default  
employee  
practice  
Time taken: 0.827 seconds, Fetched: 3 row(s)
```

**To Check the table inside the Database inside the HDFS :**

```
hive> show tables;  
OK  
bucket_sex  
people  
sex  
Time taken: 0.037 seconds, Fetched: 3 row(s)
```

**To Check the schema of the table in HDFS:**

```
hive> desc people;  
OK  
index                string  
userid               string  
firstname             string  
lastname              string  
sex                   string  
email                 string  
mob                   string  
dob                   string  
jobtitle              string  
Time taken: 0.23 seconds, Fetched: 9 row(s)
```

**To Create a table inside the HDFS :**

```
hive> create table sample(index string,userid string,firstname string,lastname string,sex  
string,email string,mob string,dob string,jobtitle string)  
> row format delimited fields terminated by ',';  
OK
```

Time taken: 1.65 seconds

### **To Load data inside the created table using a csv or tsv file :**

```
hive> load data local inpath '/home/cloudera/Training/19.07.23/people-100.csv' into table sample;
```

Loading data to table practice.sample

Table practice.sample stats: [numFiles=1, totalSize=11099]

OK

Time taken: 0.884 seconds

### **To Drop a table from the HDFS :**

```
hive> drop table sample;
```

OK

Time taken: 0.692 seconds

## **Partitioning of the table**

### **To Create a table inside the HDFS using Partitions :**

```
hive> create table sex(index string,userid string,firstname string,lastname string,email string,mob string,dob string,jobtitle string)
```

```
> partitioned by (sex string)
```

```
> row format delimited fields terminated by ',';
```

OK

Time taken: 0.084 seconds

### **Enable Permissions to accept the partitions :**

```
hive> set hive.exec.dynamic.partition = true;
```

```
hive> set hive.exec.dynamic.partition.mode = nonstrict;
```

### **Insert data inside the partitions table :**

```
hive> insert overwrite table sex
```

```
> partition(sex)
```

```
> select index,userid,firstname,lastname,email,mob,dob,jobtitle,sex from people;
```

Query ID = cloudera\_20230725000202\_0797619b-5691-4f2c-b5a8-f258f6554c64

Total jobs = 3

Launching Job 1 out of 3

Number of reduce tasks is set to 0 since there's no reduce operator

Starting Job = job\_1690266933628\_0001, Tracking URL =

[http://quickstart.cloudera:8088/proxy/application\\_1690266933628\\_0001/](http://quickstart.cloudera:8088/proxy/application_1690266933628_0001/)

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job\_1690266933628\_0001

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0

2023-07-25 00:03:11,130 Stage-1 map = 0%, reduce = 0%

2023-07-25 00:03:19,858 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 3.12 sec

MapReduce Total cumulative CPU time: 3 seconds 120 msec

Moving data to: hdfs://quickstart.cloudera:8020/user/hive/warehouse/practice.db/sex/.hive-staging\_hive\_2023-07-25\_00-02-57\_382\_4486635514177853212-1/-ext-10000

Loading data to table practice.sex partition (sex=null)

Time taken for load dynamic partitions : 362

Loading partition {sex=Male}

Loading partition {sex=Female}

Time taken for adding to write entity : 1

Partition practice.sex{sex=Female} stats: [numFiles=1, numRows=53, totalSize=5318, rawDataSize=5265]

Partition practice.sex{sex=Male} stats: [numFiles=1, numRows=47, totalSize=4741, rawDataSize=4694]

MapReduce Jobs Launched:

Stage-Stage-1: Map: 1 Cumulative CPU: 3.12 sec HDFS Read: 16078 HDFS Write: 10182 SUCCESS

Total MapReduce CPU Time Spent: 3 seconds 120 msec

OK

Time taken: 25.514 seconds

### **To Check the Partitions created:**

```
hive> dfs -ls /user/hive/warehouse/practice.db/sex;
```

Found 2 items

```
drwxrwxrwx - cloudera supergroup      0 2023-07-25 00:03
```

```
/user/hive/warehouse/practice.db/sex/sex=Female
```

```
drwxrwxrwx - cloudera supergroup      0 2023-07-25 00:03
```

```
/user/hive/warehouse/practice.db/sex/sex=Male
```

### **To Check the Contents inside the partitions created:**

```
hive> dfs -ls /user/hive/warehouse/practice.db/sex/sex=Female ;
```

Found 1 items

```
-rwxrwxrwx 1 cloudera supergroup    5318 2023-07-25 00:03
```

```
/user/hive/warehouse/practice.db/sex/sex=Female/000000_0
```

```
hive> dfs -cat /user/hive/warehouse/practice.db/sex/sex=Female/000000_0 ;
```

## Bucketing of the table

### To Create a table inside the HDFS using Buckets :

```
hive> create table bucket_sex(index string,userid string,firstname string,lastname string,sex string,email string,mob string,dob string,jobtitle string)
```

```
> clustered by (sex) into 2 buckets
```

```
> row format delimited fields terminated by ',';
```

OK

Time taken: 0.089 seconds

### Enable Permissions to accept the Bucketing :

```
hive> set hive.enforce.bucketing = true;
```

### Insert data inside the Buckets table :

```
hive> insert overwrite table bucket_sex
```

```
> select * from people;
```

Query ID = cloudera\_20230725001212\_a0ebabce-afdb-4138-9a6a-1e9bdde18400

Total jobs = 1

Launching Job 1 out of 1

Number of reduce tasks determined at compile time: 2

In order to change the average load for a reducer (in bytes):

```
set hive.exec.reducers.bytes.per.reducer=<number>
```

In order to limit the maximum number of reducers:

```
set hive.exec.reducers.max=<number>
```

In order to set a constant number of reducers:

```
set mapreduce.job.reduces=<number>
```

Starting Job = job\_1690266933628\_0002, Tracking URL =

[http://quickstart.cloudera:8088/proxy/application\\_1690266933628\\_0002/](http://quickstart.cloudera:8088/proxy/application_1690266933628_0002/)

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job\_1690266933628\_0002

Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 2

2023-07-25 00:12:44,960 Stage-1 map = 0%, reduce = 0%

2023-07-25 00:12:53,519 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 3.44 sec

2023-07-25 00:13:03,486 Stage-1 map = 100%, reduce = 50%, Cumulative CPU 7.7 sec

2023-07-25 00:13:04,580 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 11.39 sec

MapReduce Total cumulative CPU time: 11 seconds 390 msec

Ended Job = job\_1690266933628\_0002

Loading data to table practice.bucket\_sex

*Table practice.bucket\_sex stats: [numFiles=2, numRows=100, totalSize=10665, rawDataSize=10565]*

*MapReduce Jobs Launched:*

*Stage-Stage-1: Map: 1 Reduce: 2 Cumulative CPU: 11.39 sec HDFS Read: 24646 HDFS Write: 10819 SUCCESS*

*Total MapReduce CPU Time Spent: 11 seconds 390 msec*

*OK*

*Time taken: 28.894 seconds*

### **To Check the Buckets created:**

*hive> dfs -ls /user/hive/warehouse/practice.db/bucket\_sex;*

*Found 2 items*

*-rwxrwxrwx 1 cloudera supergroup 5689 2023-07-25 00:13*

*/user/hive/warehouse/practice.db/bucket\_sex/000000\_0*

*-rwxrwxrwx 1 cloudera supergroup 4976 2023-07-25 00:13*

*/user/hive/warehouse/practice.db/bucket\_sex/000001\_0*

### **To Check the Contents inside the Buckets created:**

*hive> dfs -cat /user/hive/warehouse/practice.db/bucket\_sex/000000\_0;*

### **Store the tables in different file formats :**

#### **Text file :**

*hive> create table sample\_textfile(index string,userid string,firstname string,lastname string,sex string,email string,mob string,dob string,jobtitle string)*

*> row format delimited fields terminated by ','*

*> stored as textfile;*

*OK*

*Time taken: 0.086 seconds*

### **Inserting the data inside the Textfile Format :**

*hive> insert overwrite table sample\_textfile*

*> select \* from people;*

*Query ID = cloudera\_20230725004141\_341770d5-c803-4211-97d8-60f54347e004*

*Total jobs = 3*

*Launching Job 1 out of 3*

*Number of reduce tasks is set to 0 since there's no reduce operator*

Starting Job = job\_1690266933628\_0004, Tracking URL =  
http://quickstart.cloudera:8088/proxy/application\_1690266933628\_0004/  
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job\_1690266933628\_0004  
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0  
2023-07-25 00:41:30,743 Stage-1 map = 0%, reduce = 0%  
2023-07-25 00:41:38,156 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 3.39 sec  
MapReduce Total cumulative CPU time: 3 seconds 390 msec  
Ended Job = job\_1690266933628\_0004  
Stage-4 is selected by condition resolver.  
Stage-3 is filtered out by condition resolver.  
Stage-5 is filtered out by condition resolver.  
Moving data to:  
hdfs://quickstart.cloudera:8020/user/hive/warehouse/practice.db/sample\_textfile/.hive-  
staging\_hive\_2023-07-25\_00-41-24\_810\_7553595839077842798-1/-ext-10000  
Loading data to table practice.sample\_textfile  
Table practice.sample\_textfile stats: [numFiles=1, numRows=100, totalSize=10665,  
rawDataSize=10565]  
MapReduce Jobs Launched:  
Stage-Stage-1: Map: 1 Cumulative CPU: 3.39 sec HDFS Read: 15744 HDFS Write: 10749  
SUCCESS  
Total MapReduce CPU Time Spent: 3 seconds 390 msec  
OK  
Time taken: 14.832 seconds

### **Sequence file:**

```
hive> create table sample_sequencefile(index string,userid string,firstname string,lastname  
string,sex string,email string,mob string,dob string,jobtitle string)  
> row format delimited fields terminated by ','  
> stored as sequencefile;  
OK  
Time taken: 0.09 seconds
```

### **Inserting the data inside the Sequence File Format :**

```
hive> insert overwrite table sample_sequencefile  
> select * from people;  
Query ID = cloudera_20230725004242_927084c1-5f32-45e6-9676-f405fff14e49  
Total jobs = 3  
Launching Job 1 out of 3  
Number of reduce tasks is set to 0 since there's no reduce operator
```



Starting Job = job\_1690266933628\_0005, Tracking URL =  
http://quickstart.cloudera:8088/proxy/application\_1690266933628\_0005/  
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job\_1690266933628\_0005  
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0  
2023-07-25 00:42:08,717 Stage-1 map = 0%, reduce = 0%  
2023-07-25 00:42:14,140 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.66 sec  
MapReduce Total cumulative CPU time: 1 seconds 660 msec  
Ended Job = job\_1690266933628\_0005  
Stage-4 is selected by condition resolver.  
Stage-3 is filtered out by condition resolver.  
Stage-5 is filtered out by condition resolver.  
Moving data to:  
hdfs://quickstart.cloudera:8020/user/hive/warehouse/practice.db/sample\_sequencefile/.hive-staging\_hive\_2023-07-25\_00-42-00\_714\_6855151151982995237-1/-ext-10000  
Loading data to table practice.sample\_sequencefile  
Table practice.sample\_sequencefile stats: [numFiles=1, numRows=100, totalSize=12054, rawDataSize=10565]  
MapReduce Jobs Launched:  
Stage-Stage-1: Map: 1 Cumulative CPU: 1.66 sec HDFS Read: 15882 HDFS Write: 12142  
SUCCESS  
Total MapReduce CPU Time Spent: 1 seconds 660 msec  
OK  
Time taken: 14.921 seconds

### **Parquet file:**

```
hive> create table sample_parquet(index string,userid string,firstname string,lastname  
string,sex string,email string,mob string,dob string,jobtitle string)  
> row format delimited fields terminated by ','  
> stored as parquet;  
OK  
Time taken: 0.087 seconds
```

### **Inserting the data inside the Parquet Format :**

```
hive> insert overwrite table sample_parquet  
> select * from people;  
Query ID = cloudera_20230725004343_764900c1-8b34-4367-8139-ccb81842f189  
Total jobs = 3  
Launching Job 1 out of 3  
Number of reduce tasks is set to 0 since there's no reduce operator
```

Starting Job = job\_1690266933628\_0006, Tracking URL =  
http://quickstart.cloudera:8088/proxy/application\_1690266933628\_0006/  
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job\_1690266933628\_0006  
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0  
2023-07-25 00:43:58,393 Stage-1 map = 0%, reduce = 0%  
2023-07-25 00:44:05,804 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 4.21 sec  
MapReduce Total cumulative CPU time: 4 seconds 210 msec  
Ended Job = job\_1690266933628\_0006  
Stage-4 is selected by condition resolver.  
Stage-3 is filtered out by condition resolver.  
Stage-5 is filtered out by condition resolver.  
Moving data to:  
hdfs://quickstart.cloudera:8020/user/hive/warehouse/practice.db/sample\_parquet/.hive-staging\_hive\_2023-07-25\_00-43-52\_408\_2623996529817986814-1/-ext-10000  
Loading data to table practice.sample\_parquet  
Table practice.sample\_parquet stats: [numFiles=1, numRows=100, totalSize=13693, rawDataSize=900]  
MapReduce Jobs Launched:  
Stage-Stage-1: Map: 1 Cumulative CPU: 4.21 sec HDFS Read: 16049 HDFS Write: 13774  
SUCCESS  
Total MapReduce CPU Time Spent: 4 seconds 210 msec  
OK  
Time taken: 16.001 seconds

### **Avro file:**

```
hive> create table sample_avro(index string,userid string,firstname string,lastname string,sex
string,email string,mob string,dob string,jobtitle string)
> row format delimited fields terminated by ','
> stored as avro;
OK
Time taken: 0.135 seconds
```

### **Inserting the data inside the Avro Format :**

```
hive> insert overwrite table sample_avro
> select * from people;
Query ID = cloudera_20230725015252_bffc841f-3fb5-40cf-891d-45248c591a60
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
```

Starting Job = job\_1690266933628\_0007, Tracking URL =  
http://quickstart.cloudera:8088/proxy/application\_1690266933628\_0007/  
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job\_1690266933628\_0007  
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0  
2023-07-25 01:52:49,057 Stage-1 map = 0%, reduce = 0%  
2023-07-25 01:52:56,600 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.65 sec  
MapReduce Total cumulative CPU time: 2 seconds 650 msec  
Ended Job = job\_1690266933628\_0007  
Stage-4 is selected by condition resolver.  
Stage-3 is filtered out by condition resolver.  
Stage-5 is filtered out by condition resolver.  
Moving data to:  
hdfs://quickstart.cloudera:8020/user/hive/warehouse/practice.db/sample\_avro/.hive-  
staging\_hive\_2023-07-25\_01-52-41\_291\_7381657658923414155-1/-ext-10000  
Loading data to table practice.sample\_avro  
Table practice.sample\_avro stats: [numFiles=1, numRows=100, totalSize=12212,  
rawDataSize=0]  
MapReduce Jobs Launched:  
Stage-Stage-1: Map: 1 Cumulative CPU: 2.65 sec HDFS Read: 16615 HDFS Write: 12272  
SUCCESS  
Total MapReduce CPU Time Spent: 2 seconds 650 msec  
OK  
Time taken: 16.777 seconds

### **Orc file :**

create table sample\_orc(index string,userid string,firstname string,lastname string,sex  
string,email string,mob string,dob string,jobtitle string)  
> row format delimited fields terminated by ','  
> stored as orc;  
OK  
Time taken: 0.104 seconds

### **Inserting the data inside the Orc Format :**

hive> insert overwrite table sample\_orc  
> select \* from people;  
Query ID = cloudera\_20230725003737\_a1af694d-0985-434c-840b-82eed82dcaf8  
Total jobs = 1  
Launching Job 1 out of 1  
Number of reduce tasks is set to 0 since there's no reduce operator

Starting Job = job\_1690266933628\_0003, Tracking URL =  
http://quickstart.cloudera:8088/proxy/application\_1690266933628\_0003/  
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job\_1690266933628\_0003  
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0  
2023-07-25 00:37:45,803 Stage-1 map = 0%, reduce = 0%  
2023-07-25 00:37:53,245 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.98 sec  
MapReduce Total cumulative CPU time: 1 seconds 980 msec  
Ended Job = job\_1690266933628\_0003  
Stage-4 is selected by condition resolver.  
Stage-3 is filtered out by condition resolver.  
Stage-5 is filtered out by condition resolver.  
Moving data to:  
hdfs://quickstart.cloudera:8020/user/hive/warehouse/practice.db/sample\_orc/.hive-  
staging\_hive\_2023-07-25\_00-37-36\_191\_5857662608120748751-1/-ext-10000  
Loading data to table practice.sample\_orc  
Table practice.sample\_orc stats: [numFiles=1, numRows=100, totalSize=6440,  
rawDataSize=85000]  
MapReduce Jobs Launched:  
Stage-Stage-1: Map: 1 Cumulative CPU: 1.98 sec HDFS Read: 16003 HDFS Write: 6519  
SUCCESS  
Total MapReduce CPU Time Spent: 1 seconds 980 msec  
OK  
Time taken: 18.552 seconds

**Comparing the file size of all the different file formats which has same data inserted :**

```
hive> dfs -ls /user/hive/warehouse/practice.db/sample_avro ;  
Found 1 items  
-rwxrwxrwx 1 cloudera supergroup 12212 2023-07-25 01:52  
/user/hive/warehouse/practice.db/sample_avro/000000_0
```

```
hive> dfs -ls /user/hive/warehouse/practice.db/sample_orc;  
Found 1 items  
-rwxrwxrwx 1 cloudera supergroup 6440 2023-07-25 00:37  
/user/hive/warehouse/practice.db/sample_orc/000000_0
```

```
hive> dfs -ls /user/hive/warehouse/practice.db/sample_parquet;  
Found 1 items
```

```
-rwxrwxrwx 1 cloudera supergroup 13693 2023-07-25 00:44
/user/hive/warehouse/practice.db/sample_parquet/000000_0
```

```
hive> dfs -ls /user/hive/warehouse/practice.db/sample_textfile;
Found 1 items
```

```
-rwxrwxrwx 1 cloudera supergroup 10665 2023-07-25 00:41
/user/hive/warehouse/practice.db/sample_textfile/000000_0
```

```
hive> dfs -ls /user/hive/warehouse/practice.db/sample_sequencefile;
Found 1 items
```

```
-rwxrwxrwx 1 cloudera supergroup 12054 2023-07-25 00:42
/user/hive/warehouse/practice.db/sample_sequencefile/000000_0
```

**JOIN** is a clause that is used for combining specific fields from two tables by using values common to each one. It is used to combine records from two or more tables in the database. It is more or less similar to SQL JOIN.

### Syntax

join\_table:

table\_reference JOIN table\_factor [join\_condition]

| table\_reference {LEFT|RIGHT|FULL} [OUTER] JOIN table\_reference join\_condition

| table\_reference LEFT SEMI JOIN table\_reference join\_condition

| table\_reference CROSS JOIN table\_reference [join\_condition]

### Example

We will use the following two tables in this chapter. Consider the following table named CUSTOMERS.

ID	NAME	AGE	ADDRESS	SALARY
1	Ramesh	32	Ahmedabad	2000.00
2	Khilan	25	Delhi	1500.00
3	kaushik	23	Kota	2000.00
4	Chaitali	25	Mumbai	6500.00

5	Hardik	27	Bhopal	8500.00
6	Komal	22	MP	4500.00
7	Muffy	24	Indore	10000.00

+----+-----+----+-----+-----+

Consider another table ORDERS as follows:

+----+-----+-----+-----+
OID   DATE   CUSTOMER_ID   AMOUNT
+----+-----+-----+-----+
102   2009-10-08 00:00:00   3   3000
100   2009-10-08 00:00:00   3   1500
101   2009-11-20 00:00:00   2   1560
103   2008-05-20 00:00:00   4   2060

+----+-----+-----+-----+

There are different types of joins given as follows:

- JOIN
- LEFT OUTER JOIN
- RIGHT OUTER JOIN
- FULL OUTER JOIN

## JOIN

JOIN clause is used to combine and retrieve the records from multiple tables. JOIN is same as OUTER JOIN in SQL. A JOIN condition is to be raised using the primary keys and foreign keys of the tables.

The following query executes JOIN on the CUSTOMER and ORDER tables, and retrieves the records:

```
hive> SELECT c.ID, c.NAME, c.AGE, o.AMOUNT
FROM CUSTOMERS c JOIN ORDERS o ON (c.ID = o.CUSTOMER_ID);
```

On successful execution of the query, we get to see the following response:

+----+-----+----+-----+

ID	NAME	AGE	AMOUNT
----	------	-----	--------

3	kaushik	23	1500
---	---------	----	------

2	Khilan	25	1560
---	--------	----	------

4	Chaitali	25	2060
---	----------	----	------

4	Chaitali	25	2060
---	----------	----	------

4	Chaitali	25	2060
---	----------	----	------

### LEFT OUTER JOIN

The HiveQL LEFT OUTER JOIN returns all the rows from the left table, even if there are no matches in the right table. This means, if the ON clause matches 0 zero records in the right table, the JOIN still returns a row in the result, but with NULL in each column from the right table.

A LEFT JOIN returns all the values from the left table, plus the matched values from the right table, or NULL in case of no matching JOIN predicate.

The following query demonstrates LEFT OUTER JOIN between CUSTOMER and ORDER tables:

```
hive> SELECT c.ID, c.NAME, o.AMOUNT, o.DATE FROM CUSTOMERS c
```

```
LEFT OUTER JOIN ORDERS o ON (c.ID = o.CUSTOMER_ID);
```

On successful execution of the query, we get to see the following response:

1	Ramesh	NULL	NULL
---	--------	------	------

2	Khilan	1560	2009-11-20 00:00:00
---	--------	------	---------------------

3	kaushik	3000	2009-10-08 00:00:00
---	---------	------	---------------------

3	kaushik	1500	2009-10-08 00:00:00
---	---------	------	---------------------

4	Chaitali	2060	2008-05-20 00:00:00
---	----------	------	---------------------

5	Hardik	NULL	NULL
---	--------	------	------

6	Komal	NULL	NULL
---	-------	------	------

7	Muffy	NULL	NULL
---	-------	------	------

7	Muffy	NULL	NULL
---	-------	------	------

7	Muffy	NULL	NULL
---	-------	------	------

7	Muffy	NULL	NULL
---	-------	------	------

7	Muffy	NULL	NULL
---	-------	------	------

## RIGHT OUTER JOIN

The HiveQL RIGHT OUTER JOIN returns all the rows from the right table, even if there are no matches in the left table. If the ON clause matches 0 zero records in the left table, the JOIN still returns a row in the result, but with NULL in each column from the left table.

A RIGHT JOIN returns all the values from the right table, plus the matched values from the left table, or NULL in case of no matching join predicate.

The following query demonstrates RIGHT OUTER JOIN between the CUSTOMER and ORDER tables.

```
hive> SELECT c.ID, c.NAME, o.AMOUNT, o.DATE FROM CUSTOMERS c RIGHT  
OUTER JOIN ORDERS o ON c. ID = o. CUSTOMERID;
```

On successful execution of the query, we get to see the following response:

ID	NAME	AMOUNT	DATE
3	kaushik	1500	2009-10-08 00:00:00
2	Khilan	1560	2009-11-20 00:00:00
4	Chaitali	2060	2008-05-20 00:00:00

## FULL OUTER JOIN

The HiveQL FULL OUTER JOIN combines the records of both the left and the right outer tables that fulfil the JOIN condition. The joined table contains either all the records from both the tables, or fills in NULL values for missing matches on either side.

The following query demonstrates FULL OUTER JOIN between CUSTOMER and ORDER tables:

```
hive> SELECT c.ID, c.NAME, o.AMOUNT, o.DATE FROM CUSTOMERS c  
FULL OUTER JOIN ORDERS o ON (c.ID = o.CUSTOMER_ID);
```

ID	NAME	AMOUNT	DATE
1	Ramesh	NULL	NULL
2	Khilan	1560	2009-11-20 00:00:00
3	kaushik	3000	2009-10-08 00:00:00



3	kaushik	1500	2009-10-08 00:00:00	
4	Chaitali	2060	2008-05-20 00:00:00	
5	Hardik	NULL	NULL	
6	Komal	NULL	NULL	
7	Muffy	NULL	NULL	
3	kaushik	3000	2009-10-08 00:00:00	
3	kaushik	1500	2009-10-08 00:00:00	
2	Khilan	1560	2009-11-20 00:00:00	
4	Chaitali	2060	2008-05-20 00:00:00	

+-----+-----+-----+-----+